

The background of the slide features a series of concentric circles in various shades of blue and light blue, creating a ripple effect that radiates from the center. The text is centered within this design.

# **DATA SCIENCE FOR SUPPLY CHAIN FORECASTING**

by Diego Beteta

# Intro

## Pasado

Empresas tecnológicas proporcionaron software de pronóstico estadístico, basados en suavización exponencial, que permitieron a las empresas utilizarlas como el pilar de sus procesos S&OP.

## Presente

Debido al aumento de la potencia computacional, grandes conjuntos de datos, mejores modelos de pronóstico y disponibilidad de herramientas gratuitas, las empresas pueden ser cada vez más competitivas.

Con modelos de **Machine Learning (ML)** construidos en Python, es posible aportar a cualquier negocio más valor que cualquier software de pronóstico disponible en el mercado.

## Futuro

Los 'demand planners' tendrán que aprender a trabajar con modelos de pronóstico avanzados basados en ML y podrán agregarles valor a medida que comprendan las deficiencias de ML.



“Si Excel es una navaja suiza, Python es un ejército completo de máquinas de construcción que esperan instrucciones de cualquier científico de datos.”

# Excel vs Python

Es necesario un cambio de paradigma para pasar aproximaciones manuales realizadas en Excel a modelos potentes automatizados en Python.

Python permite realizar cálculos en grandes conjuntos de datos de forma rápida y automatizada. Viene con muchas bibliotecas:

- Análisis de datos (pandas)
- Cálculos científicos (Numpy/SciPy)
- Machine Learning (scikit-learn)

Beneficios de aplicar ‘Data Science’ en el pronóstico de la demanda:

- Escalabilidad
- Automatización
- Rentabilidad

# CONTENT

+

•

Application of Data Science in Demand Management

Old-School Statistics vs Machine Learning

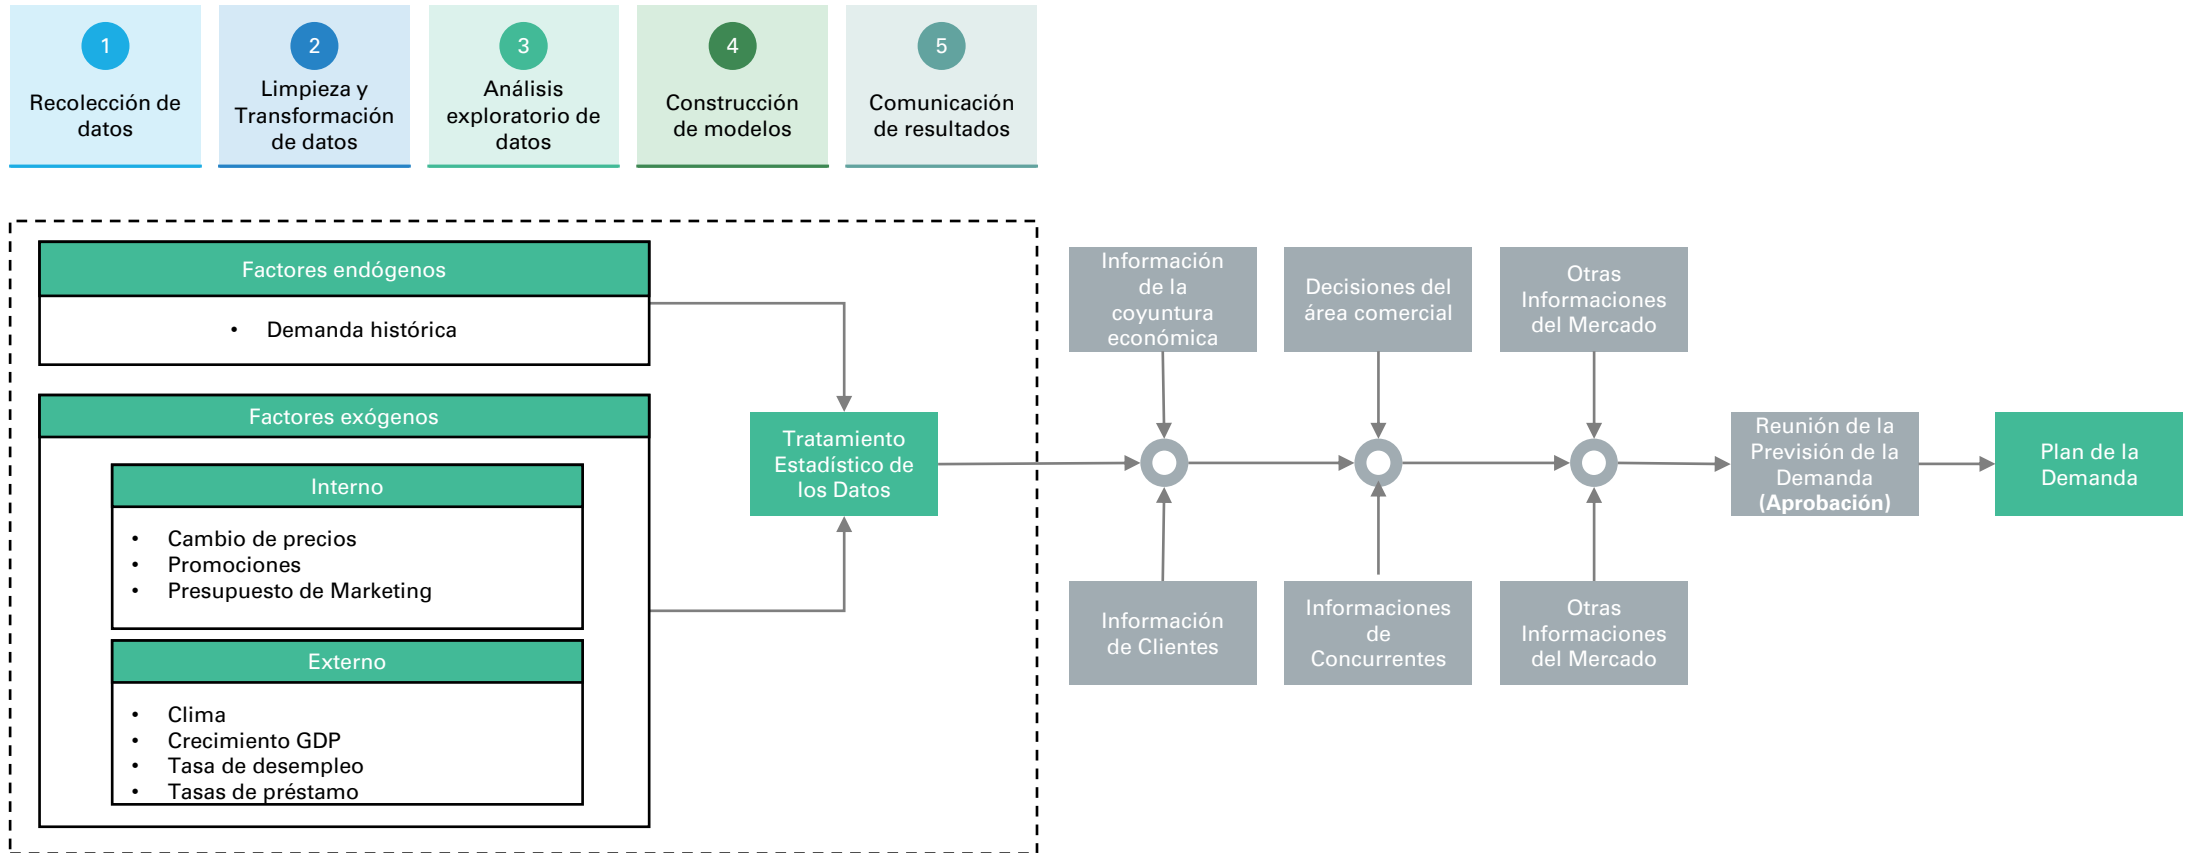
Forecast KPI

Hands-on

Wrap-up

○

# Application of Data Science in Demand Management



# Old-School Statistics vs Machine Learning

## Statistical Forecasting

- Permite ver y entender los patrones de la demanda (variación aleatoria, tendencia y estacionalidad)
- Sólo permite procesar un producto a la vez
- No permite evaluar la relación de datos externos con datos históricos de la demanda

### Modelos de pronóstico

1. Moving Average
2. Exponential Smoothing
3. Double Exponential Smoothing
4. Double Exponential Smoothing with Damped Trend
5. Triple Exponential Smoothing
6. Tripple Additive Exponential Smoothing

## Machine Learning

- No proporciona ninguna explicación ni comprensión de los diferentes patrones de demanda
- Sólo se enfoca en obtener la respuesta correcta
- Permite procesar una gran cantidad de productos a la vez
- Permite evaluar la relación de datos externos con datos históricos de la demanda

### Modelos de pronóstico

1. Linear Regression
2. Tree
3. Forest
4. Extremely Randomized Trees
5. Adaptative Boosting
6. Extreme Gradient Boosting

# Forecast KPI

Definir la calidad del pronóstico mediante la evaluación de sus KPI es importante porque nos permite:

- Dimensionar la confiabilidad de cada técnica de pronóstico
- Mejorar la calidad de cada técnica mediante la optimización de sus parámetros
- Comparar diferentes técnicas optimizadas y definir las más adecuada para nuestros productos.
- Supervisar constantemente las técnicas utilizadas

## Adaptabilidad

No existe un modelo de pronóstico perfecto que pueda vencer a cualquier otro modelo para cualquier tipo de negocio.

Adaptar los modelos de pronóstico a su conjunto de datos de demanda permitirá lograr un mejor nivel de precisión que mediante el uso de herramientas de caja negra.

Los modelos de Machine Learning deben adaptarse a los patrones de demanda de sus productos.






Figure 2.6: Demand and forecasts.

Table 2.1: KPI comparison.

	Forecast #1	Forecast #2	Forecast #3
Bias	-3.9	-1.9	0.1
MAPE	64%	109%	180%
MAE	4.4	4.1	4.8
RMSE	7.1	6.2	5.9

# Forecast KPI

BIAS Average Error	MAPE Mean Absolute Percentage Error	MAE Mean Absolute Error	RMSE Root Mean Square Error
<ul style="list-style-type: none"><li>• El sesgo promedio de un pronóstico se define como su error promedio.</li><li>• Nos indica si, en promedio, subestimamos o sobrestimamos la demanda.</li><li>• Un valor relativamente bajo (%) indica que la técnica no tiene sesgo, no sobrestima ni subestima consistentemente la demanda.</li></ul>	<ul style="list-style-type: none"><li>• Es especialmente útil cuando los valores de la demanda real son muy grandes.</li><li>• MAPE divide cada error individualmente por la demanda, por lo que está sesgado: un pronóstico extremadamente bajo sólo puede resultar en un error máximo de 100%, mientras que cualquier pronóstico demasiado alto no se limitará a un porcentaje de error específico (&gt;100%)</li><li>• Debido a esto, la optimización de MAPE dará como resultado un pronóstico que muy probablemente subestime la demanda. ¡Evita usar MAPE!</li></ul>	<ul style="list-style-type: none"><li>• Mide la precisión del pronóstico al promediar las magnitudes de los errores absolutos del pronóstico de cada periodo de tiempo.</li><li>• Nos indica qué tan lejos, en promedio, están nuestros pronósticos de la demanda mediana.</li><li>• Un valor relativamente alto (%) indican que se tiene una pésima calidad de pronóstico.</li></ul>	<ul style="list-style-type: none"><li>• RMSE le da más importancia a los errores más grandes, mientras que MAE le da la misma importancia a cada error.</li><li>• Nos indica qué tan lejos, en promedio, están nuestros pronósticos de la demanda promedio.</li><li>• Un gran error es suficiente para obtener un RMSE muy malo (alto %).</li></ul>



# HANDS-ON

## IMPORT, TRANSFORM AND FILTER DATA

## DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

dataset.csv X

dataset.csv

1	Year,Month,Product,Demand
2	2007,1,Product 1,2884
3	2007,1,Product 2,2521
4	2007,1,Product 3,1029
5	2007,1,Product 4,870
6	2007,1,Product 5,693
7	2007,1,Product 6,665
8	2007,1,Product 7,622
9	2007,1,Product 8,599
10	2007,1,Product 9,423
11	2007,1,Product 10,362
12	2007,1,Product 11,352
13	2007,1,Product 12,263
14	2007,1,Product 13,258
15	2007,1,Product 14,191
16	2007,1,Product 15,169
17	2007,1,Product 16,168
18	2007,1,Product 17,136
19	2007,1,Product 18,127
20	2007,1,Product 19,97
21	2007,1,Product 20,55
22	2007,1,Product 21,33
23	2007,1,Product 22,26
24	2007,1,Product 23,26
25	2007,1,Product 24,22
26	2007,1,Product 25,20
27	2007,1,Product 26,16
28	2007,1,Product 27,15
29	2007,1,Product 28,14
30	2007,1,Product 29,9
31	2007,1,Product 30,4
32	2007,1,Product 31,4
33	2007,1,Product 32,3
34	2007,1,Product 33,2
35	2007,1,Product 34,2
36	2007,1,Product 35,2
37	2007,1,Product 36,1
38	2007,1,Product 37,1
39	2007,1,Product 38,1
40	2007,2,Product 1,1885
41	2007,2,Product 2,1517
42	2007,2,Product 4,686
43	2007,2,Product 3,621
44	2007,2,Product 5,570
45	2007,2,Product 7,551
46	2007,2,Product 8,498



```
# Import dataset to dataframe
def import_data():
    data = pd.read_csv('dataset.csv')
    return data

data = import_data()
display(data)
print('There are a total of', data['Product'].nunique(),
      | 'unique products in the dataset')
```

[166] ✓ 0.2s

...

	Year	Month	Product	Demand
0	2007	1	Product 1	2884
1	2007	1	Product 2	2521
2	2007	1	Product 3	1029
3	2007	1	Product 4	870
4	2007	1	Product 5	693
...	...	...	...	...
4372	2017	1	Product 34	3
4373	2017	1	Product 46	2
4374	2017	1	Product 48	1
4375	2017	1	Product 36	1
4376	2017	1	Product 37	1

4377 rows × 4 columns

There are a total of 66 unique products in the dataset



```
# To filter the product
product = 'Product 40'
dataset = data[data['Product'] == product]
display(dataset)
```

[177] ✓ 0.6s

...

	Year	Month	Product	Demand
74	2007	2	Product 40	2
105	2007	3	Product 40	14
145	2007	4	Product 40	7
184	2007	5	Product 40	4
222	2007	6	Product 40	5
...	...	...	...	...
4229	2016	9	Product 40	26
4263	2016	10	Product 40	14
4299	2016	11	Product 40	29
4335	2016	12	Product 40	27
4363	2017	1	Product 40	91

117 rows × 4 columns

# HANDS-ON

## 1 MOVING AVERAGE

```
d = dataset['Demand'].tolist() # Convert dataframe column to a list
df = moving_average(d, extra_periods=4, n=3) # Call our new function
```

... Moving Average Forecast - Product 40

Period	Demand	Forecast	Error
0	2.0	NaN	NaN
1	14.0	NaN	NaN
2	7.0	NaN	NaN
3	4.0	7.666667	-3.666667
4	5.0	8.333333	-3.333333
...	...	...	...
116	91.0	23.333333	67.666667
117	NaN	49.000000	NaN
118	NaN	49.000000	NaN
119	NaN	49.000000	NaN
120	NaN	49.000000	NaN

[121 rows x 3 columns]

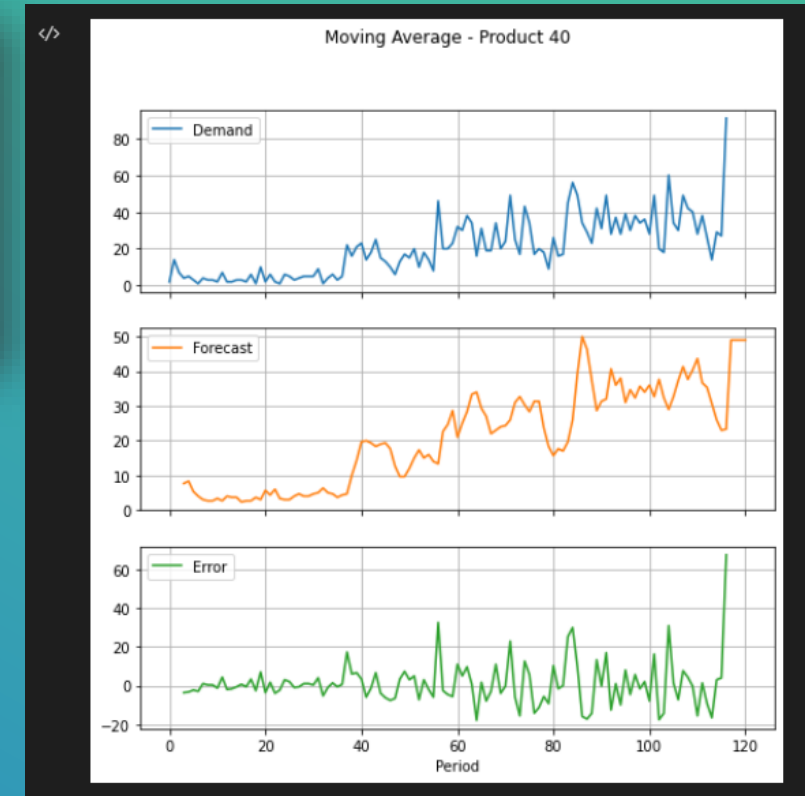
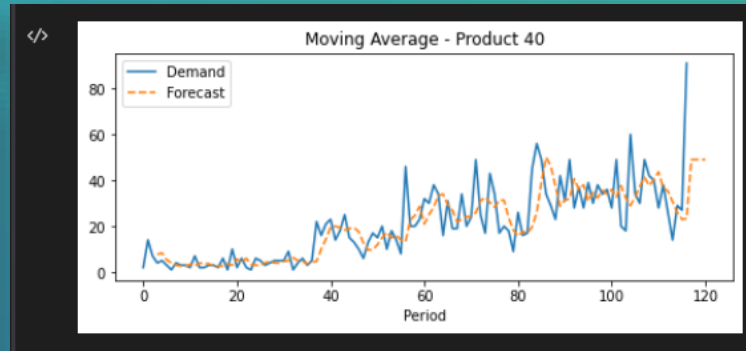
Moving Average KPI - Product 40

Bias: 0.89, 4.27%

MAPE: 50.79%

MAE: 7.28, 34.87%

RMSE: 11.45, 54.88%



# HANDS-ON

## 2 SIMPLE EXPONENTIAL SMOOTHING

### DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

```
d = dataset['Demand'].tolist() # Convert dataframe column to list
df = simple_exp_smooth(d, extra_periods=4, alpha=0.4)
```

... Simple Exponential Smoothing Forecast - Product 40

Period	Demand	Forecast	Error
0	2.0	NaN	NaN
1	14.0	2.000000	12.000000
2	7.0	6.800000	0.200000
3	4.0	6.880000	-2.880000
4	5.0	5.728000	-0.728000
...	...	...	...
116	91.0	26.756136	64.243864
117	NaN	52.453682	NaN
118	NaN	52.453682	NaN
119	NaN	52.453682	NaN
120	NaN	52.453682	NaN

[121 rows x 3 columns]

Simple Exponential Smoothing KPI - Product 40


Bias: 1.09, 5.26%

MAPE: 50.94%


MAE: 7.23, 34.93%

RMSE: 11.27, 54.45%

</>



</>



# HANDS-ON

## 3 DOUBLE EXPONENTIAL SMOOTHING

### DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

```
d = dataset['Demand'].tolist() # Convert dataframe column to a list
df = double_exp_smooth(d, extra_periods=4, alpha=0.4,
                        beta=0.4)
```





```
... Double Exponential Smoothing Forecast - Product 40
```

Period	Demand	Forecast	Level	Trend	Error
0	2.0	NaN	2.000000	12.000000	NaN
1	14.0	14.000000	14.000000	12.000000	0.000000
2	7.0	26.000000	18.400000	8.960000	-19.000000
3	4.0	27.360000	18.016000	5.222400	-23.360000
4	5.0	23.238400	15.943040	2.304256	-18.238400
...	...	...	...	...	...
116	91.0	20.161385	48.496831	9.235834	70.838615
117	NaN	57.732665	57.732665	9.235834	NaN
118	NaN	66.968500	66.968500	9.235834	NaN
119	NaN	76.204334	76.204334	9.235834	NaN
120	NaN	85.440169	85.440169	9.235834	NaN

```
[121 rows x 5 columns]
Double Exponential Smoothing KPI - Product 40
Bias: -0.15, -0.72%
MAPE: 76.64%
MAE: 8.66, 41.86%
RMSE: 13.07, 63.19%
```

**Level:** es el valor medio en torno al cual varía la demanda a lo largo del tiempo.

**Trend:** es la variación promedio del nivel de la serie temporal entre dos periodos consecutivos.



**Relación tendencia-error:** la tendencia disminuye cuando el error es positivo y la tendencia aumenta cuando el error es negativo. La intuición es que nuestro modelo aprende de sus errores.

Si el modelo subestima la última demanda, aumentará la tendencia. Si sobrestima la última demanda, disminuirá la tendencia.

# HANDS-ON

## 4 DOUBLE EXPONENTIAL SMOOTHING WITH DAMPED TREND

### DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

```
d = dataset['Demand'].tolist() # Convert dataframe column to a list
df = double_exp_smooth_damped(d, extra_periods=4, alpha=0.4,
                               beta=0.4, phi=0.9)
```

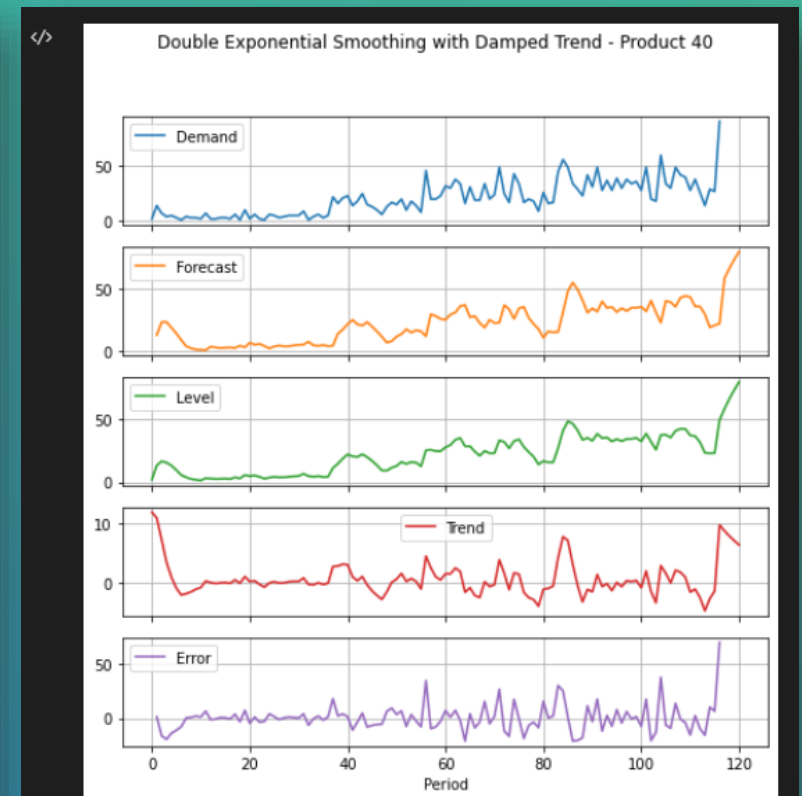

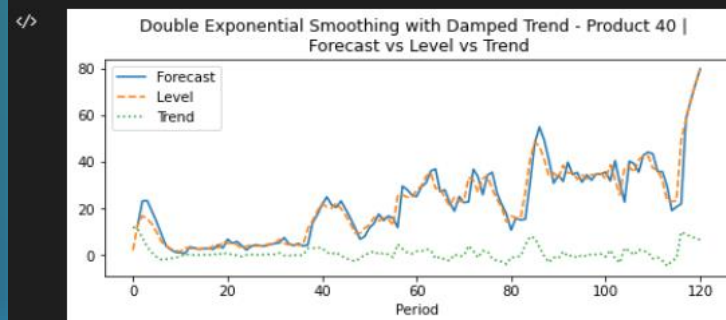

```
*** Double Exponential Smoothing with Damped Trend Forecast - Product 40
```

Period	Demand	Forecast	Level	Trend	Error
0	2.0	NaN	2.000000	12.000000	NaN
1	14.0	12.800000	13.280000	10.992000	1.200000
2	7.0	23.172800	16.703680	7.305152	-16.172800
3	4.0	23.278317	15.566990	3.490106	-19.278317
4	5.0	18.708086	13.224851	0.947802	-13.708086
...	...	...	...	...	...
116	91.0	21.910738	49.546443	9.797628	69.089262
117	NaN	58.364308	58.364308	8.817865	NaN
118	NaN	66.300387	66.300387	7.936079	NaN
119	NaN	73.442858	73.442858	7.142471	NaN
120	NaN	79.871082	79.871082	6.428224	NaN

```
[121 rows x 5 columns]
Double Exponential Smoothing with Damped Trend KPI - Product 40
Bias: 0.13, 0.63%
MAPE: 68.53%
MAE: 8.23, 39.76%
RMSE: 12.51, 60.48%
```

### Factor de amortiguamiento (phi):

Reducirá exponencialmente la tendencia a lo largo del tiempo. Este nuevo modelo se olvida de la tendencia con el paso del tiempo. O que el modelo recuerde solo una fracción (phi) de la tendencia estimada anterior.





# HANDS-ON

## 5 TRIPLE EXPONENTIAL SMOOTHING






### DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

```
d = dataset['Demand'].tolist() # Convert dataframe column to a list
df = triple_exp_smooth_mul(d, slen=12, extra_periods=4, alpha=0.4,
                           beta=0.4, phi=0.9, gamma=0.2)
```

```
... Triple Exponential Smoothing Forecast - Product 40
```

Period	Demand	Forecast	Level	Trend	Season	Error
0	2.0	NaN	1.924187	10.336338	1.039400	NaN
1	14.0	12.819718	11.640345	9.468086	1.141876	1.180282
2	7.0	21.743069	14.693320	6.333956	1.078438	-14.743069
3	4.0	22.192582	13.706649	3.025668	1.088198	-18.192582
4	5.0	12.426984	12.502056	1.152023	0.756371	-7.426984
...	...	...	...	...	...	...
116	91.0	36.021133	46.968460	7.473446	1.375607	54.978867
117	NaN	45.102712	53.694561	6.726101	0.839987	NaN
118	NaN	44.534721	59.748052	6.053491	0.745375	NaN
119	NaN	66.744360	65.196194	5.448142	1.023746	NaN
120	NaN	71.466505	70.099522	4.903328	1.019501	NaN

[121 rows x 6 columns]  
Triple Exponential Smoothing KPI - Product 40  
Bias: -0.24, -1.17%  
MAPE: 70.45%  
MAE: 7.96, 38.47%  
RMSE: 11.99, 57.94%



# HANDS-ON

## 6 TRIPLE ADDITIVE EXPONENTIAL SMOOTHING

### DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

```
d = dataset['Demand'].tolist() # Convert dataframe column to a list
df = triple_exp_smooth_add(d, slen=12, extra_periods=4, alpha=0.4,
                           beta=0.4, phi=0.9, gamma=0.2)
```

```
''' Triple Additive Exponential Smoothing Forecast - Product 40
Demand Forecast Level Trend Season Error
Period
0      2.0      NaN    1.192593  9.900000  0.807407      NaN
1     14.0    13.010000  10.498593  9.068400  2.907407    0.990000
2      7.0    20.267560  13.353129  6.038750  1.607407   -13.267560
3      4.0    20.595411  12.149839  2.779610  1.807407   -16.595411
4      5.0     9.658895  12.787930  1.756225  -4.992593   -4.658895
...
116    91.0    35.505386  49.340805  8.712197  15.021781   55.494614
117     NaN    54.623637  57.181782  7.840977  -2.558145      NaN
118     NaN    59.611517  64.238661  7.056880  -4.627144      NaN
119     NaN    77.161796  70.589853  6.351192  6.571944      NaN
120     NaN    78.760716  76.305925  5.716072  2.454791      NaN
```

[121 rows x 6 columns]

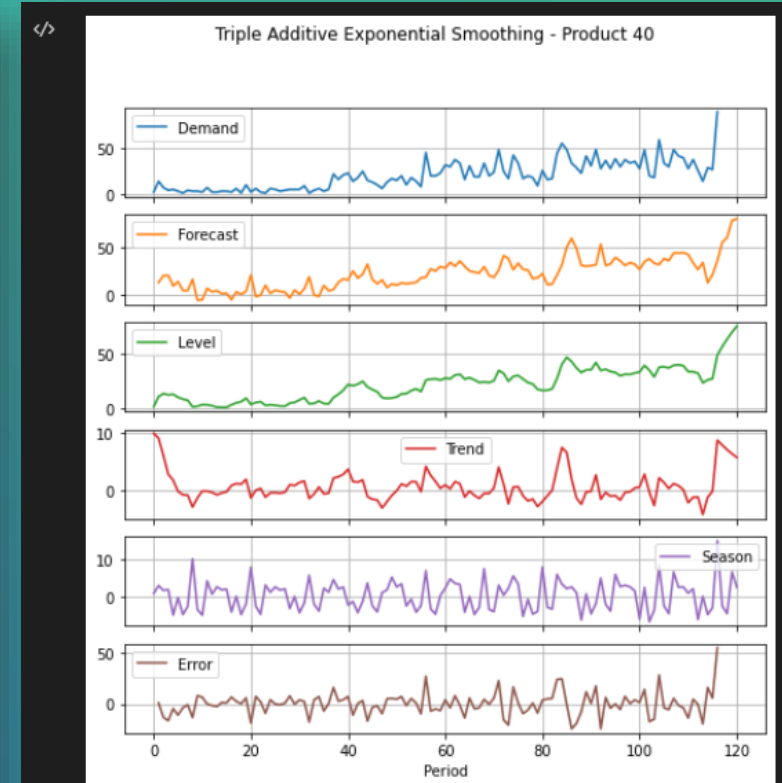




Triple Additive Exponential Smoothing KPI - Product 40

Bias: 0.18, 0.85%

MAPE: 94.63%

MAE: 8.02, 38.74%

RMSE: 11.50, 55.58%



# HANDS-ON

## DEFINE THE BEST MODEL FOR OUR PRODUCT OPTIMAL RMSE

## DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

```
def exp_smooth_opti_rmse(d, extra_periods=6, slen=12):
    params = [] # contains all the different parameter sets
    dfs = [] # contains all the DataFrames returned by the different models
    KPIs = [] # contains the results of each model

    for alpha in [0.05, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6]:
        df = simple_exp_smooth(d, extra_periods=extra_periods, alpha=alpha)
        params.append(f'Simple Exponential Smoothing, alpha: {alpha}')
        dfs.append(df)
        RMSE = np.sqrt((df['Error']**2).mean())
        KPIs.append(RMSE)

    for beta in [0.05, 0.1, 0.2, 0.3, 0.4]:
        df = double_exp_smooth(d, extra_periods=extra_periods, alpha=alpha, beta=beta)
        params.append(f'Double Exponential Smoothing, alpha: {alpha}, beta: {beta}')
        dfs.append(df)
        RMSE = np.sqrt((df['Error']**2).mean())
        KPIs.append(RMSE)

    for phi in [0.7, 0.75, 0.8, 0.85, 0.9, 0.95, 1.00]: # To forget the trend over time
        df = double_exp_smooth_damped(d, extra_periods=extra_periods, alpha=alpha, beta=beta, phi=phi)
        params.append(f'Double Exponential Smoothing with Damped Trend, alpha: {alpha}, beta: {beta}, phi: {phi}')
        dfs.append(df)
        RMSE = np.sqrt((df['Error']**2).mean())
        KPIs.append(RMSE)

    for gamma in [0.05, 0.1, 0.2, 0.3]:
        df = triple_exp_smooth_mul(d, slen=slen, extra_periods=extra_periods, alpha=alpha, beta=beta, phi=phi, gamma=gamma)
        params.append(f'Triple Exponential Smoothing, alpha: {alpha}, beta: {beta}, phi: {phi}, gamma: {gamma}')
        dfs.append(df)
        RMSE = np.sqrt((df['Error']**2).mean())
        KPIs.append(RMSE)

        df = triple_exp_smooth_add(d, slen=slen, extra_periods=extra_periods, alpha=alpha, beta=beta, phi=phi, gamma=gamma)
        params.append(f'Triple Additive Exponential Smoothing, alpha: {alpha}, beta: {beta}, phi: {phi}, gamma: {gamma}')
        dfs.append(df)
        RMSE = np.sqrt((df['Error']**2).mean())
        KPIs.append(RMSE)

    mini = np.argmin(KPIs)
    # np.argmin() returns the location of the minimum value in an array.
    # Remember that, in Python, the index (location) of the first element in an array is 0.
    print(f'Best solution found: {params[mini]}. RMSE: ', round(KPIs[mini], 2))
    return dfs[mini]
```

```
# Forecast
d = dataset['Demand'].tolist() # Convert dataframe column to a list
df = exp_smooth_opti_rmse(d, extra_periods=6, slen=12)
df.index.name = 'Period'
print(df)
```

```
... Best solution found: Triple Exponential Smoothing, alpha: 0.3, beta: 0.05, phi: 0.7, gamma: 0.05. RMSE: 10.01
```

	Demand	Forecast	Level	Trend	Season	Error
Period						
0	2.0	NaN	1.924187	10.336338	1.039400	NaN
1	14.0	10.459155	10.089894	7.281950	1.141876	3.540845
2	7.0	16.378525	12.578342	4.966919	1.078438	-9.378525
3	4.0	17.471222	12.341370	3.291153	1.088198	-13.471222
4	5.0	11.077185	12.234778	2.183287	0.756371	-6.077185
...	...	...	...	...	...	...
118	NaN	31.881295	41.689297	0.234740	0.764736	NaN
119	NaN	48.024932	41.853615	0.164318	1.147450	NaN
120	NaN	44.270222	41.968638	0.115023	1.054841	NaN
121	NaN	48.782631	42.049153	0.080516	1.160133	NaN
122	NaN	45.687254	42.105514	0.056361	1.085066	NaN

[123 rows x 6 columns]



# HANDS-ON

## DEFINE THE BEST MODEL FOR OUR PRODUCT OPTIMAL RMSE






### DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

\*\*\* Best solution found: Triple Exponential Smoothing, alpha: 0.3, beta: 0.05, phi: 0.7, gamma: 0.05. RMSE: 10.01

\*\*\* Triple Exponential Smoothing Forecast - Product 40

Period	Demand	Forecast	Level	Trend	Season	Error
0	2.0	NaN	1.924187	10.336338	1.039400	NaN
1	14.0	10.459155	10.089894	7.281950	1.141876	3.540845
2	7.0	16.378525	12.578342	4.966919	1.078438	-9.378525
3	4.0	17.471222	12.341370	3.291153	1.088198	-13.471222
4	5.0	11.077185	12.234778	2.183287	0.756371	-6.077185
...	...	...	...	...	...	...
116	91.0	42.237217	41.119215	0.479061	1.424462	48.762783
117	NaN	35.363314	41.454557	0.335343	0.853062	NaN
118	NaN	31.881295	41.689297	0.234740	0.764736	NaN
119	NaN	48.024932	41.853615	0.164318	1.147450	NaN
120	NaN	44.270222	41.968638	0.115023	1.054841	NaN

[121 rows x 6 columns]  
Triple Exponential Smoothing KPI - Product 40  
Bias: 0.42, 2.03%  
MAPE: 67.47%  
MAE: 6.67, 32.23%  
RMSE: 10.01, 48.40%



# HANDS-ON

## DEFINE THE BEST MODEL FOR OUR PRODUCT OPTIMAL MAE

## DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

```
def exp_smooth_opti_mae(d, extra_periods=6, slen=12):
    params = [] # contains all the different parameter sets
    dfs = [] # contains all the DataFrames returned by the different models
    KPIs = [] # contains the results of each model

    for alpha in [0.05, 0.1, 0.2, 0.3, 0.4, 0.5, 0.6]:
        df = simple_exp_smooth(d, extra_periods=extra_periods, alpha=alpha)
        params.append(f'Simple Exponential Smoothing, alpha: {alpha}')
        dfs.append(df)
        MAE = df['Error'].abs().mean()
        KPIs.append(MAE)

    for beta in [0.05, 0.1, 0.2, 0.3, 0.4]:
        df = double_exp_smooth(d, extra_periods=extra_periods, alpha=alpha, beta=beta)
        params.append(f'Double Exponential Smoothing, alpha: {alpha}, beta: {beta}')
        dfs.append(df)
        MAE = df['Error'].abs().mean()
        KPIs.append(MAE)

    for phi in [0.7, 0.75, 0.8, 0.85, 0.9, 0.95, 1.00]: # To forget the trend over time
        df = double_exp_smooth_damped(d, extra_periods=extra_periods, alpha=alpha, beta=beta, phi=phi)
        params.append(f'Double Exponential Smoothing with Damped Trend, alpha: {alpha}, beta: {beta}, phi: {phi}')
        dfs.append(df)
        MAE = df['Error'].abs().mean()
        KPIs.append(MAE)

    for gamma in [0.05, 0.1, 0.2, 0.3]:
        df = triple_exp_smooth_mul(d, slen=slen, extra_periods=extra_periods, alpha=alpha, beta=beta, phi=phi, gamma=gamma)
        params.append(f'Triple Exponential Smoothing, alpha: {alpha}, beta: {beta}, phi: {phi}, gamma: {gamma}')
        dfs.append(df)
        MAE = df['Error'].abs().mean()
        KPIs.append(MAE)

        df = triple_exp_smooth_add(d, slen=slen, extra_periods=extra_periods, alpha=alpha, beta=beta, phi=phi, gamma=gamma)
        params.append(f'Triple Additive Exponential Smoothing, alpha: {alpha}, beta: {beta}, gamma: {gamma}')
        dfs.append(df)
        MAE = df['Error'].abs().mean()
        KPIs.append(MAE)

    mini = np.argmin(KPIs)
    # np.argmin() returns the location of the minimum value in an array.
    # Remember that, in Python, the index (location) of the first element in an array is 0.
    print(f'Best solution found: {params[mini]}. MAE: ', round(KPIs[mini], 2))
    return dfs[mini]
```

```
# Forecast
d = dataset['Demand'].tolist() # Convert dataframe column to a list
df = exp_smooth_opti_mae(d, extra_periods=6, slen=12)
df.index.name = 'Period'
print(df)
```

... Best solution found: Triple Exponential Smoothing, alpha: 0.4, beta: 0.05, phi: 0.7, gamma: 0.05. MAE: 6.59

Period	Demand	Forecast	Level	Trend	Season	Error
0	2.0	NaN	1.924187	10.336338	1.039400	NaN
1	14.0	10.459155	10.399985	7.297455	1.141876	3.540845
2	7.0	16.724643	11.901268	4.927872	1.078438	-9.724643
3	4.0	16.704688	10.680788	3.216011	1.088198	-12.704688
4	5.0	9.781384	10.403403	2.124778	0.756371	-4.781384
...	...	...	...	...	...	...
118	NaN	34.669020	45.594792	0.334064	0.760372	NaN
119	NaN	52.226430	45.828637	0.233845	1.139603	NaN
120	NaN	48.388994	45.992329	0.163691	1.052110	NaN
121	NaN	53.054483	46.106912	0.114584	1.150684	NaN
122	NaN	49.982213	46.187121	0.080209	1.082168	NaN

[123 rows x 6 columns]

# HANDS-ON

## DEFINE THE BEST MODEL FOR OUR PRODUCT OPTIMAL MAE

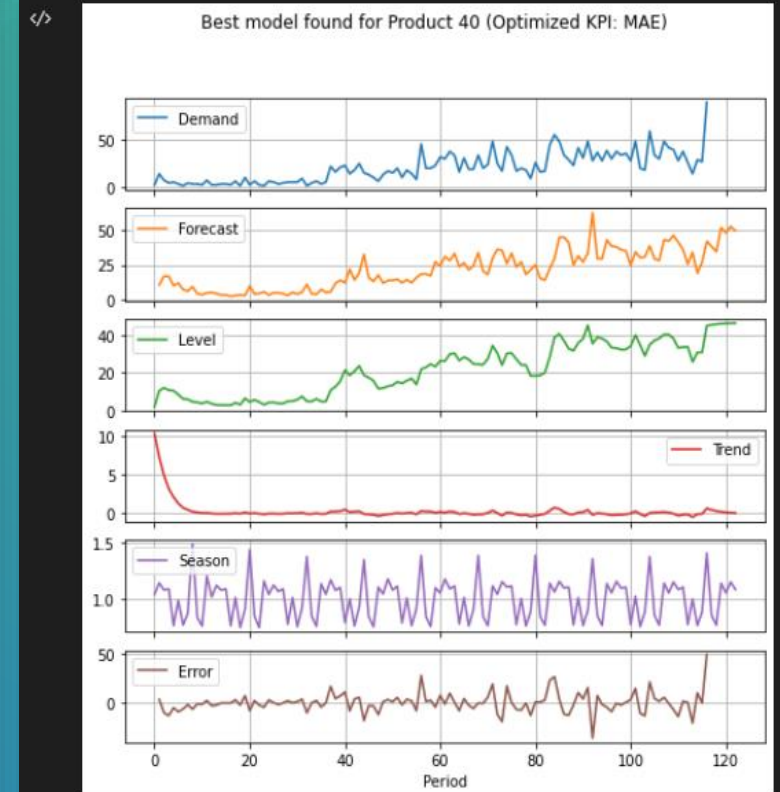




## DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

\*\*\* Best solution found: Triple Exponential Smoothing, alpha: 0.4, beta: 0.05, phi: 0.7, gamma: 0.05. MAE: 6.59

\*\*\* Triple Exponential Smoothing Forecast - Product 40

Period	Demand	Forecast	Level	Trend	Season	Error
0	2.0	NaN	1.924187	10.336338	1.039400	NaN
1	14.0	10.459155	10.399985	7.297455	1.141876	3.540845
2	7.0	16.724643	11.901268	4.927872	1.078438	-9.724643
3	4.0	16.704688	10.680788	3.216011	1.088198	-12.704688
4	5.0	9.781384	10.403403	2.124778	0.756371	-4.781384
...	...	...	...	...	...	...
116	91.0	42.262323	44.783494	0.681763	1.411670	48.737677
117	NaN	38.460452	45.260728	0.477234	0.849753	NaN
118	NaN	34.669020	45.594792	0.334064	0.760372	NaN
119	NaN	52.226430	45.828637	0.233845	1.139603	NaN
120	NaN	48.388994	45.992329	0.163691	1.052110	NaN

[121 rows x 6 columns]  
Triple Exponential Smoothing KPI - Product 40  
Bias: 0.38, 1.86%  
MAPE: 64.03%  
MAE: 6.59, 31.87%  
RMSE: 10.10, 48.79%



# HANDS-ON

## DEAL WITH OUTLIERS (AUTOMATICALLY)

La detección de valores atípicos es un asunto serio. Estos valores atípicos aparecen todo el tiempo en las cadenas de suministro modernas.

En su mayoría se deben a errores o una demanda excepcional. Si puede detectar valores atípicos y suavizarlos, hará un mejor pronóstico.

Marcar valores atípicos manualmente es un proceso que requiere mucho tiempo, es propenso a errores y no es gratificante; pocos planificadores de la demanda se tomarán el tiempo necesario para revisarlos.

Por lo tanto, cuanto mayor sea el conjunto de datos, más importante es automatizar esta detección y limpieza.

## DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

Con el método de detección más inteligente, analizando la desviación estándar del error de pronóstico, podremos marcar los valores atípicos con mucha más precisión y reemplazarlos con una cantidad plausible.

Se tomará dos desviaciones estándar como límites para este método.

```
# Import norm to compute probabilities
from scipy.stats import norm

# Compute the error mean and error standard deviation
m = df['Error'].mean()
s = df['Error'].std()

# Get the percentage of the normal distribution function of each error with norm.cdf
prob = norm.cdf(df['Error'], m, s) # Normal cumulative distribution function

# Define probabilities ranges based on 2 standard deviation
outliers = (prob > 0.99) | (prob < 0.01)

# Re-compute the error mean and standard deviation, but this time, exclude the outliers. '~' useful for exclude.
m2 = df.loc[~outliers, 'Error'].mean()
s2 = df.loc[~outliers, 'Error'].std()

# Update the lower and upper acceptable thresholds based on these new m2 and s2 values
limit_high = norm.ppf(0.99, m2, s2) + df['Forecast'] # Percent point function
limit_low = norm.ppf(0.01, m2, s2) + df['Forecast'] # Percent point function

# Forecast with updated outliers values
df['Updated'] = df['Demand'].clip(lower=limit_low, upper=limit_high)
display(df)

# Replace the updated values in the main dataset
updated_values = df.loc[df['Updated'].notnull(), 'Updated']
dataset['Demand'] = dataset['Demand'].replace(dataset['Demand'].tolist(), updated_values)
print('Dataset with Demand updated')
print(dataset)
```

	Demand	Forecast	Level	Trend	Season	Error	Updated
Period							
113	14.0	34.239264	25.610577	-0.461364	0.997103	-20.239264	15.899767
114	29.0	18.919780	30.676793	-0.053496	0.758041	10.080220	29.000000
115	27.0	26.928939	30.671687	-0.035830	0.878970	0.071061	27.000000
116	91.0	42.262323	44.783494	0.681763	1.411670	48.737677	60.189113
117	NaN	38.460452	45.260728	0.477234	0.849753	NaN	NaN
118	NaN	34.669020	45.594792	0.334064	0.760372	NaN	NaN
119	NaN	52.226430	45.828637	0.233845	1.139603	NaN	NaN
120	NaN	48.388994	45.992329	0.163691	1.052110	NaN	NaN
121	NaN	53.054483	46.106912	0.114584	1.150684	NaN	NaN
122	NaN	49.982213	46.187121	0.080209	1.082168	NaN	NaN

Dataset with Demand updated					
	Year	Month	Product	Demand	
74	2007	2	Product 40	2.000000	
105	2007	3	Product 40	15.899767	
145	2007	4	Product 40	7.000000	
184	2007	5	Product 40	4.000000	
222	2007	6	Product 40	5.000000	
...	...	...	...	...	
4229	2016	9	Product 40	26.000000	
4263	2016	10	Product 40	15.899767	
4299	2016	11	Product 40	29.000000	
4335	2016	12	Product 40	27.000000	
4363	2017	1	Product 40	60.189113	



# HANDS-ON

## DEFINE THE BEST MODEL FOR OUR PRODUCT OPTIMAL RMSE (WITHOUT OUTLIERS)


### DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

```
# Forecast
d = dataset['Demand'].tolist() # Convert dataframe column to a list
df = exp_smooth_opti_rmse(d, extra_periods=6, slen=12)
df.index.name = 'Period'
print(df)
```





```
*** Triple Exponential Smoothing Forecast - Product 40
```

Period	Demand	Forecast	Level	Trend	Season	Error
0	2.000000	NaN	1.963651	11.568077	1.018511	NaN
1	15.899767	11.822023	11.102432	8.149710	1.174999	4.077744
2	7.000000	18.501136	13.672791	5.548075	1.100784	-11.501136
3	4.000000	19.500754	13.369865	3.674324	1.110746	-15.500754
4	5.000000	12.307832	13.102220	2.430043	0.772043	-7.307832
...	...	...	...	...	...	...
116	60.189113	37.178948	35.699665	0.222749	1.257361	23.010165
117	NaN	31.061854	35.855589	0.155924	0.866304	NaN
118	NaN	27.967106	35.964736	0.109147	0.777626	NaN
119	NaN	41.218501	36.041139	0.076403	1.143651	NaN
120	NaN	37.405204	36.094621	0.053482	1.036310	NaN

[121 rows x 6 columns]  
Triple Exponential Smoothing KPI - Product 40  
Bias: 0.17, 0.84%  
MAPE: 67.38%  
MAE: 6.28, 31.02%  
RMSE: 8.58, 42.36%



\*\*\* Best solution found: Triple Exponential Smoothing, alpha: 0.3, beta: 0.05, phi: 0.7, gamma: 0.05. RMSE: 8.58



# HANDS-ON

## DEFINE THE BEST MODEL FOR OUR PRODUCT OPTIMAL MAE (WITHOUT OUTLIERS)


### DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

```
# Forecast
d = dataset['Demand'].tolist() # Convert dataframe column to a list
df = exp_smooth_opti_mae(d, extra_periods=6, slen=12)
df.index.name = 'Period'
print(df)
```

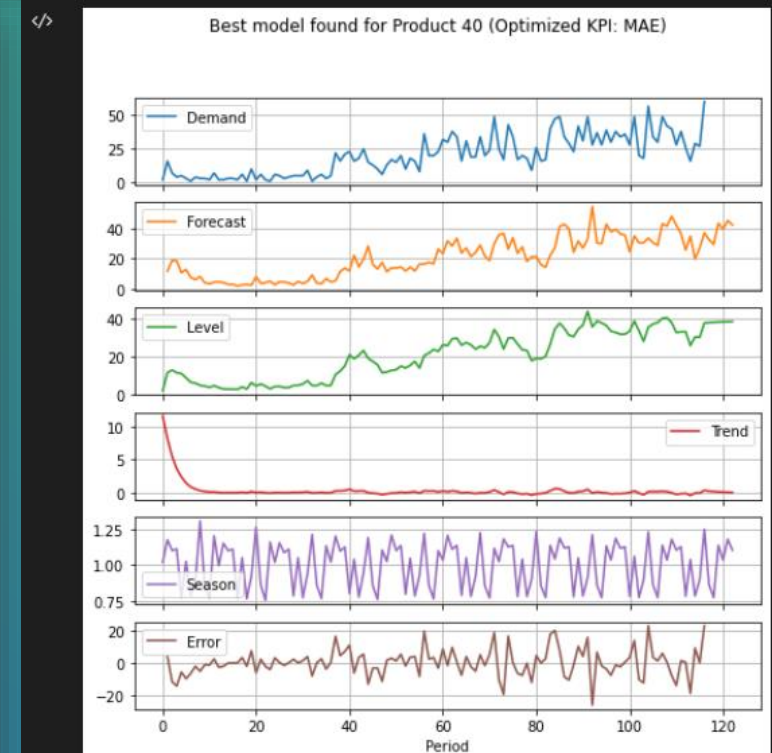



```
*** Triple Exponential Smoothing Forecast - Product 40
```

Period	Demand	Forecast	Level	Trend	Season	Error
0	2.000000	NaN	1.963651	11.568077	1.018511	NaN
1	15.899767	11.822023	11.449474	8.167062	1.174999	4.077744
2	7.000000	18.896525	12.843491	5.500797	1.100784	-11.896525
3	4.000000	18.542853	11.456902	3.588701	1.110746	-14.542853
4	5.000000	10.784668	10.971924	2.362237	0.772043	-5.784668
...	...	...	...	...	...	...
116	60.189113	37.055260	37.663412	0.342569	1.247970	23.133853
117	NaN	32.675162	37.903211	0.239799	0.862068	NaN
118	NaN	29.434172	38.071070	0.167859	0.773138	NaN
119	NaN	43.411720	38.188571	0.117501	1.136773	NaN
120	NaN	39.573321	38.270822	0.082251	1.034034	NaN

[121 rows x 6 columns]  
Triple Exponential Smoothing KPI - Product 40  
Bias: 0.16, 0.78%  
MAPE: 63.06%  
MAE: 6.21, 30.63%  
RMSE: 8.60, 42.42%



\*\*\* Best solution found: Triple Exponential Smoothing, alpha: 0.4, beta: 0.05, phi: 0.7, gamma: 0.05. MAE: 6.21



# WRAP-UP

## THE ANALYSIS

### DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

De los 65 únicos productos del conjunto de datos, se escogió aleatoriamente el producto # 40.

Detectando y reemplazando con cantidades plausibles a las demandas atípicas del producto, podemos mejorar nuestro pronóstico.

Si nuestro objetivo es ajustar nuestro pronóstico a la demanda promedio, debemos enfocarnos en reducir lo máximo posible el KPI: RMSE (%)

Para el producto #40, con demandas atípicas ajustadas, el algoritmo de optimización sugiere que la técnica de pronóstico estadístico 'Triple Exponential Smoothing' es la más adecuada con los parámetros:

- alpha: 0.3
- beta: 0.05
- phi: 0.7
- gamma: 0.05

Dando un RMSE de 8.58 unidades (42.36%)

El gráfico 'Demand vs Season' indica que el producto #40 es estacional.

```
... Best solution found: Triple Exponential Smoothing, alpha: 0.3, beta: 0.05, phi: 0.7, gamma: 0.05. RMSE: 8.58
```

```
... Best solution found: Triple Exponential Smoothing, alpha: 0.4, beta: 0.05, phi: 0.7, gamma: 0.05. MAE: 6.21
```

Si nuestro objetivo es ajusta nuestro pronóstico a la demanda mediana, debemos enfocarnos en reducir lo máximo posible el KPI: MAE (%)

Para el producto #40, con demandas atípicas ajustadas, el algoritmo de optimización sugiere que la técnica de pronóstico estadístico 'Triple Exponential Smoothing' es la más adecuada con los parámetros:

- alpha: 0.4
- beta: 0.05
- phi: 0.7
- gamma: 0.05

Dando un MAE de 6.21 unidades (30.63%)

El gráfico 'Demand vs Season' indica que el producto #40 es estacional.

Para el producto #40, el pronóstico de los próximos 04 meses son 31.06, 27.97, 47.21 y 37.41, respectivamente.

Sin embargo, ¿cómo podemos pronosticar la demanda del siguiente periodo de los más de 60 productos en nuestro conjunto de datos?

¿Cómo podemos relacionar la demanda histórica de nuestros productos con sus respectivos cambios de precios, clima, crecimiento del GDP o tas de desempleo?

**¡Esto será posible aplicando modelos de Machine Learning!**

# MACHINE LEARNING

## IMPORT , TRANSFORM DATA AND BUILT MAIN FUNCTIONS

## DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

```
dataset.csv X
dataset.csv
1 Year,Month,Product,Demand
2 2007,1,Product 1,2884
3 2007,1,Product 2,2521
4 2007,1,Product 3,1029
5 2007,1,Product 4,870
6 2007,1,Product 5,693
7 2007,1,Product 6,665
8 2007,1,Product 7,622
9 2007,1,Product 8,599
10 2007,1,Product 9,423
11 2007,1,Product 10,362
12 2007,1,Product 11,352
13 2007,1,Product 12,263
14 2007,1,Product 13,258
15 2007,1,Product 14,191
16 2007,1,Product 15,169
17 2007,1,Product 16,168
18 2007,1,Product 17,136
19 2007,1,Product 18,127
20 2007,1,Product 19,97
21 2007,1,Product 20,55
22 2007,1,Product 21,33
23 2007,1,Product 22,26
24 2007,1,Product 23,26
25 2007,1,Product 24,22
26 2007,1,Product 25,20
27 2007,1,Product 26,16
28 2007,1,Product 27,15
29 2007,1,Product 28,14
30 2007,1,Product 29,9
31 2007,1,Product 30,4
32 2007,1,Product 31,4
33 2007,1,Product 32,3
34 2007,1,Product 33,2
35 2007,1,Product 34,2
36 2007,1,Product 35,2
37 2007,1,Product 36,1
38 2007,1,Product 37,1
39 2007,1,Product 38,1
40 2007,2,Product 1,1885
41 2007,2,Product 2,1517
42 2007,2,Product 4,686
43 2007,2,Product 3,621
44 2007,2,Product 5,570
45 2007,2,Product 7,551
46 2007,2,Product 8,498
```



```
# Import dataset
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt

def import_data():
    data = pd.read_csv('dataset.csv')
    # z.fill(x) adds leading zeros to a string until it reaches a length of x.
    data['Period'] = data['Year'].astype(str) + '-' + data['Month'].astype(str).str.zfill(2)
    df = pd.pivot_table(data=data, values='Demand', index='Product', columns='Period',
                        aggfunc='sum', fill_value=0)
    return df

# Build an appropriate dataframe for our machine learning
def datasets(df, x_len=12, y_len=1, test_loops=12):
    D = df.values
    rows, periods = D.shape

    # Training set creation
    loops = periods + 1 - x_len - y_len
    train = []
    for col in range(loops):
        train.append(D[:, col:x_len+y_len])
    train = np.vstack(train)
    X_train, Y_train = np.split(train, [-y_len], axis=1)

    # Test set creation
    # np.split(array, indices, axis): cuts an array at specific indices along an axis
    if test_loops > 0:
        X_test, Y_test = np.split(X_train, [-rows*test_loops], axis=0)
        Y_train, Y_test = np.split(Y_train, [-rows*test_loops], axis=0)

    # No test set
    # X_test is used to generate the future forecast
    else:
        X_test = D[:, -x_len:]
        Y_test = np.full((X_test.shape[0], y_len), np.nan)

    # Formatting required for scikit-learn
    # array.ravel(): reduces the dimension of a Numpy array to 1D.
    if y_len == 1:
        Y_train = Y_train.ravel()
        Y_test = Y_test.ravel()

    return X_train, Y_train, X_test, Y_test
```



```
# We can now easily call our new functions
df = import_data()
X_train, Y_train, X_test, Y_test = datasets(df, x_len=12, y_len=1, test_loops=12)
print(df)

# Note that we set test_loops as 12 periods.
# That means that we will test our algorithm over 12 different loops.
# We will predict 12 times the following period [y_len=1] based on the last 12 periods [x_len=12]

# We obtain the datasets we need to training our machine learning algorithm (X_train and Y_train)
# and the datasets we need to test it (X_test and Y_test)
```

Period	2007-01	2007-02	2007-03	2007-04	2007-05	2007-06	2007-07 \
Product							
Product 1	2884	1885	1833	1300	1866	1620	1901
Product 10	362	410	387	387	422	421	469
Product 11	352	335	365	360	431	477	403
Product 12	263	247	239	179	223	277	281
Product 13	258	264	333	347	420	262	296
...	...	...	...	...	...	...	...
Product 65	0	0	0	0	0	0	0
Product 66	0	0	0	0	0	0	0
Product 7	622	551	578	534	771	683	685
Product 8	599	498	682	556	630	498	562
Product 9	423	356	399	351	520	624	401

```
#KPI Function
def kpi_ML(Y_train, Y_train_pred, Y_test, Y_test_pred, name=''):
    df = pd.DataFrame(columns = ['MAE', 'RMSE', 'Bias'], index=['Train', 'Test'])
    df.index.name = name
    df.loc['Train', 'MAE'] = 100*np.mean(abs(Y_train - Y_train_pred))/np.mean(Y_train)
    df.loc['Train', 'RMSE'] = 100*np.sqrt(np.mean((Y_train - Y_train_pred)**2))/np.mean(Y_train)
    df.loc['Train', 'Bias'] = 100*np.mean((Y_train - Y_train_pred))/np.mean(Y_train)
    df.loc['Test', 'MAE'] = 100*np.mean(abs(Y_test - Y_test_pred))/np.mean(Y_test)
    df.loc['Test', 'RMSE'] = 100*np.sqrt(np.mean((Y_test - Y_test_pred)**2))/np.mean(Y_test)
    df.loc['Test', 'Bias'] = 100*np.mean((Y_test - Y_test_pred))/np.mean(Y_test)
    df = df.astype(float).round(1) # Round number for display
    print(df)
```



# MACHINE LEARNING

## 1 LINER REGRESSION

```
from sklearn.linear_model import LinearRegression
reg = LinearRegression()

# Fit the linear regression to the training data
reg = reg.fit(X_train, Y_train)

# Create predictions for the training and test sets
Y_train_pred = reg.predict(X_train)
Y_test_pred = reg.predict(X_test)

# Measure its accuracy
kpi ML(Y_train, Y_train_pred, Y_test, Y_test_pred, name='Linear Regression')
```

✓ 1m 4.9s

	MAE	RMSE	Bias
Linear Regression			
Train	17.9	44.2	-0.0
Test	17.8	44.0	1.6

## DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

```
# Linear Regression Forecast for the next month
df = import_data()
X_train, Y_train, X_test, Y_test = datasets(df, x_len=12, y_len=1, test_loops=0)

reg = LinearRegression()
reg = reg.fit(X_train, Y_train)
forecast = pd.DataFrame(data=reg.predict(X_test), index=df.index)
print('Linear Regression Forecast for the next period')
display(forecast)
```

[6] ✓ 5.9s

... Linear Regression Forecast for the next period

```
</>
0
Product
Product 1  1457.102925
Product 10  794.012107
Product 11  1265.040299
Product 12  158.945172
Product 13  292.543366
...
Product 65  1.114935
Product 66  4.216304
Product 7   259.025417
Product 8   646.563604
Product 9   123.639567
```

66 rows × 1 columns

# MACHINE LEARNING

## 2 TREE

## DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

```
from sklearn.tree import DecisionTreeRegressor
tree = DecisionTreeRegressor(max_depth=5, min_samples_split=15, min_samples_leaf=5)

# Fit the tree to the training data
tree = tree.fit(X_train, Y_train)

# Create predictions for the training and test sets
Y_train_pred = tree.predict(X_train)
Y_test_pred = tree.predict(X_test)

# Measure its accuracy
kpi ML(Y_train, Y_train_pred, Y_test, Y_test_pred, name='Tree')
```

✓ 23.6s

MAE RMSE Bias

Tree

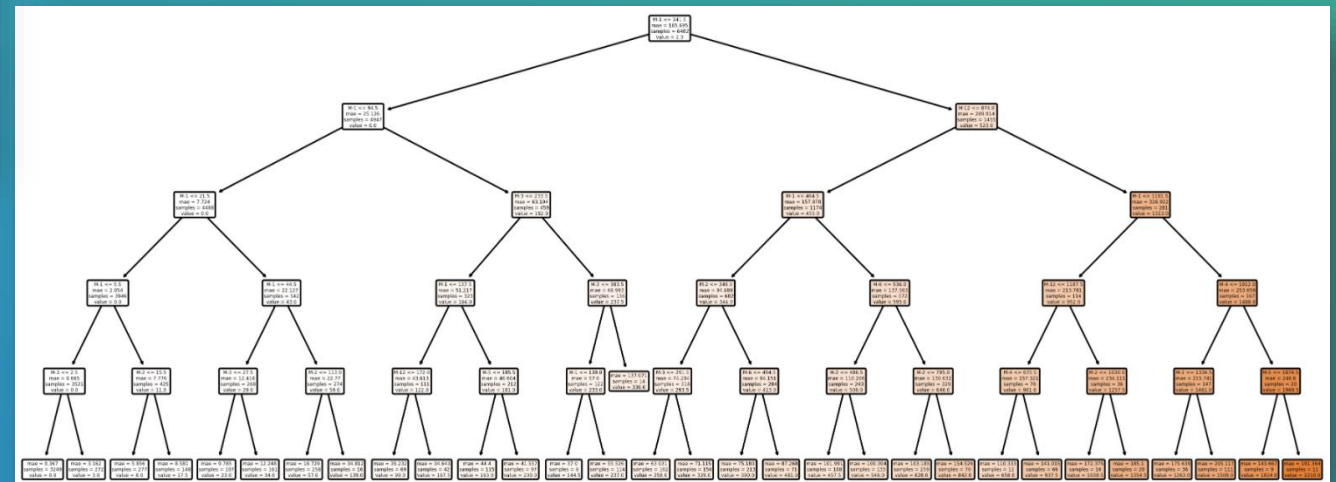
Train 18.1 43.7 -0.0

Test 21.1 53.0 3.2

```
# Plotting our tree
from sklearn.tree import plot_tree
import matplotlib.pyplot as plt
fig = plt.figure(figsize=(15,6), dpi=300)
ax = fig.gca()

plot_tree(tree, fontsize=3, feature_names=[f'M{x-12}' for x in range(12)], rounded=True, filled=True, ax=ax)
fig.savefig('Regression Tree.PNG')
```

✓ 1m 8.4s



# MACHINE LEARNING

## 2 OPTIMAL TREE

## DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

```
▷ # k-fold Cross-validation to avoid overfitting (cv parameter)
# Random Search to efficiently find a very good parameter set among different possibilities

max_depth = list(range(5,11)) + [None]
min_samples_split = range(5,20)
min_samples_leaf = range(2,20)

param_dist = {'max_depth': max_depth,
              'min_samples_split': min_samples_split,
              'min_samples_leaf': min_samples_leaf}

# Test all these different parameter sets against our training set
from sklearn.model_selection import RandomizedSearchCV
tree = DecisionTreeRegressor()
tree_cv = RandomizedSearchCV(tree, param_dist, n_jobs=-1, cv=10, verbose=1, n_iter=100, scoring='neg_mean_absolute_error')

# Fit the tree to the training data
tree_cv.fit(X_train, Y_train)
print('Tuned Regression Tree Parameters:', tree_cv.best_params_)

# Create predictions for the training and test sets
Y_train_pred = tree_cv.predict(X_train)
Y_test_pred = tree_cv.predict(X_test)

# Measure its accuracy
kpi_ML(Y_train, Y_train_pred, Y_test, Y_test_pred, name='Tree Optimization')
```

[6] ✓ 1m 18.1s

```
... Fitting 10 folds for each of 100 candidates, totalling 1000 fits
Tuned Regression Tree Parameters: {'min_samples_split': 8, 'min_samples_leaf': 14, 'max_depth': 6}

MAE  RMSE  Bias
Tree Optimization
Train  17.1  42.0   0.0
Test   19.3  48.1   2.9
```

```
# DecisionTreeRegressor Optimized Forecast for the next month
df = import_data()
X_train, Y_train, X_test, Y_test = datasets(df, x_len=12, y_len=1, test_loops=0)

# Set with the optimized parameters found
# Tuned Regression Tree Parameters: {'min_samples_split': 17, 'min_samples_leaf': 18, 'max_depth': 7}
tree = DecisionTreeRegressor(max_depth=7, min_samples_split=17, min_samples_leaf=18)
tree = tree.fit(X_train, Y_train)
forecast = pd.DataFrame(data=tree.predict(X_test), index=df.index)
print('DecisionTreeRegressor Optimized Forecast for the next period')
display(forecast)
```

[15] ✓ 1.2s

```
... DecisionTreeRegressor Optimized Forecast for the next period
```

	0
Product	
Product 1	1602.172414
Product 10	937.971429
Product 11	1216.391304
Product 12	193.126866
Product 13	265.626866
...	...
Product 65	0.162784
Product 66	10.811024
Product 7	339.521739
Product 8	591.570000
Product 9	103.406250

66 rows x 1 columns

# MACHINE LEARNING

## 3 FOREST VS OPTIMAL FOREST

### DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

```
from sklearn.ensemble import RandomForestRegressor
forest = RandomForestRegressor(bootstrap=True, max_samples=0.95, max_features=11, min_samples_leaf=18, max_depth=7)

# Fit the forest to the training data
forest.fit(X_train, Y_train)

# Create predictions for the training and test sets
Y_train_pred = forest.predict(X_train)
Y_test_pred = forest.predict(X_test)

# Measure its accuracy
kpi ML(Y_train, Y_train_pred, Y_test, Y_test_pred, name='Forest')
```

✓ 8.6s

	MAE	RMSE	Bias
Forest			
Train	15.7	40.3	0.0
Test	18.3	47.3	3.5

```
# k-fold Cross-validation to avoid overfitting (cv parameter)
# Random Search to efficiently find a very good parameter set among different possibilities

max_depth = list(range(5,11)) + [None]
min_samples_split = range(5,20)
min_samples_leaf = range(2,15)

# Bootstrap means that each tree will receive a random selection from the initial training dataset
# max_samples means to limit the amount of data given to each tree
# max_features means to limit the maximum number of features that the algorithm can choose from each node and
# the features are chosen randomly each time, we will obtain different trees at each fitting.
bootstrap = [True] #We force bootstrap
max_samples = [.7, .8, .9, .95, 1]
max_features = range(3,8)

param_dist = {'max_depth': max_depth,
              'min_samples_split': min_samples_split,
              'min_samples_leaf': min_samples_leaf,
              'bootstrap': bootstrap,
              'max_samples': max_samples,
              'max_features': max_features}

# Test all these different parameter sets against our training set
from sklearn.model_selection import RandomizedSearchCV
forest = RandomForestRegressor(n_jobs=-1, n_estimators=30)
forest_cv = RandomizedSearchCV(forest, param_dist, cv=6, n_jobs=-1, verbose=2, n_iter=400, scoring='neg_mean_absolute_error')

# Fit the forest to the training data
forest_cv.fit(X_train, Y_train)
print('Tuned Forest Parameters:', forest_cv.best_params_)

# Create predictions for the training and test sets
Y_train_pred = forest_cv.predict(X_train)
Y_test_pred = forest_cv.predict(X_test)

# Measure the accuracy
kpi ML(Y_train, Y_train_pred, Y_test, Y_test_pred, name='Forest Optimization')
```

✓ 23m 42.5s

Fitting 6 folds for each of 400 candidates, totalling 2400 fits

Tuned Forest Parameters: {'min\_samples\_split': 7, 'min\_samples\_leaf': 4, 'max\_samples': 0.9, 'max\_features': 6, 'max\_depth': 10, 'bootstrap': True}

	MAE	RMSE	Bias
Forest Optimization			
Train	12.1	30.8	0.1
Test	17.7	45.8	2.8

# MACHINE LEARNING

## 3 OPTIMAL FORESTX200

## DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

```
▷ # In order to get the best out of our forest, let's run a forest with
# our new optimal parameters and n_estimators=200.
# We can easily allow ourselves 200 trees due to the dataset limited size
forest = RandomForestRegressor(n_jobs=-1, n_estimators=200, **forest_cv.best_params_)

# Fit the forest to the training data
forest = forest.fit(X_train, Y_train)

# Create predictions for the training and test sets
Y_train_pred = forest.predict(X_train)
Y_test_pred = forest.predict(X_test)

# Measure the accuracy
kpi ML(Y_train, Y_train_pred, Y_test, Y_test_pred, name='Forestx200')
```

[10] ✓ 12.9s

...	MAE	RMSE	Bias
Forestx200			
Train	11.9	30.4	-0.0
Test	17.5	45.4	2.3

```
# Forest Optimized x 200 Forecast for the next month
df = import_data()
X_train, Y_train, X_test, Y_test = datasets(df, x_len=12, y_len=1, test_loops=0)

forest200 = RandomForestRegressor(n_jobs=-1, n_estimators=200, **forest_cv.best_params_)
forest200 = forest200.fit(X_train, Y_train)
forecast = pd.DataFrame(data=forest200.predict(X_test), index=df.index)
print('Forest Optimized x 200 Forecast for the next period')
display(forecast)
```

[6] ✓ 11.6s

... Forest Optimized x 200 Forecast for the next period

</>	0
Product	
Product 1	1525.167175
Product 10	804.875529
Product 11	1078.683406
Product 12	169.169076
Product 13	305.843355
...	...
Product 65	0.086483
Product 66	9.101335
Product 7	295.069647
Product 8	692.907534
Product 9	121.680588

66 rows × 1 columns

# MACHINE LEARNING

## 4 EXTREMELY RANDOMIZED TREES VS OPTIMAL EXTREMELY RANDOMIZED TREES

DATA SCIENCE FOR SUPPLY CHAIN  
FORECASTING

```
from sklearn.ensemble import ExtraTreesRegressor
ETR = ExtraTreesRegressor(n_jobs=-1, n_estimators=200, min_samples_split=15, min_samples_leaf=4, max_samples=0.95,
                           max_features=4, max_depth=8, bootstrap=True)

# Fit the ExtraTreesRegressor to the training data
ETR.fit(X_train, Y_train)

# Create predictions for the training and test sets
Y_train_pred = ETR.predict(X_train)
Y_test_pred = ETR.predict(X_test)

# Measure the accuracy
kpi ML(Y_train, Y_train_pred, Y_test, Y_test_pred, name='ETR')
```

[8] ✓ 7.1s

	MAE	RMSE	Bias
ETR			
Train	17.8	43.9	0.1
Test	18.8	47.5	3.3

```
# k-fold Cross-validation to avoid overfitting (cv parameter)
# Random Search to efficiently find a very good parameter set among different possibilities

max_depth = list(range(6, 13)) + [None]
min_samples_split = range(7,16)
min_samples_leaf = range(2,13)

# Bootstrap means that each tree will receive a random selection from the initial training dataset
# max_samples means to limit the amount of data ggiven to each tree
# max_features means to limit the maximum number of features that the algorithm can choose from each node and
# the features are choosen randomly each time, we will obtain different trees at each fitting.
bootstrap = [True] # We force bootstrap
max_samples = [.7, .8, .9, .95, 1]
max_features = range(5,13)

param_dist = {'max_depth': max_depth,
              'min_samples_split': min_samples_split,
              'min_samples_leaf': min_samples_leaf,
              'bootstrap': bootstrap,
              'max_samples': max_samples,
              'max_features': max_features}

# Test all these different parameter sets against our training set
ETR = ExtraTreesRegressor(n_jobs=-1, n_estimators=30)
ETR_cv = RandomizedSearchCV(ETR, param_dist, cv=5, verbose=2, n_jobs=-1, n_iter=400, scoring='neg_mean_absolute_error')

# Fit the forest to the training data
ETR_cv.fit(X_train, Y_train)
print('Tuned Forest Parameters', ETR_cv.best_params_)

# Create predictions for the training and test sets
Y_train_pred = ETR_cv.predict(X_train)
Y_test_pred = ETR_cv.predict(X_test)

# Measure the accuracy
kpi ML(Y_train, Y_train_pred, Y_test, Y_test_pred, name='ETR Optimization')
```

[10] ✓ 8m 50.9s

... Fitting 5 folds for each of 400 candidates, totalling 2000 fits

Tuned Forest Parameters {'min\_samples\_split': 11, 'min\_samples\_leaf': 2, 'max\_samples': 0.9, 'max\_features': 10, 'max\_depth': 11, 'bootstrap': True}

	MAE	RMSE	Bias
ETR optimized			
Train	14.4	36.2	0.1
Test	17.8	45.8	2.6

# MACHINE LEARNING

## 4 OPTIMAL EXTREMELY RANDOMIZED TREESX200

DATA SCIENCE FOR SUPPLY CHAIN  
FORECASTING

```
# In order to get the best out of our ETR, let's run a ETR with
# our new optimal parameters and n_estimators=200.
# We can easily allow ourselves 200 trees due to the dataset limited size
ETR = ExtraTreesRegressor(n_jobs=-1, n_estimators=200, **ETR_cv.best_params_)

# Fit the ETR to the training data
ETR = ETR.fit(X_train, Y_train)

# Create predictions for the training and test sets
Y_train_pred = ETR.predict(X_train)
Y_test_pred = ETR.predict(X_test)

# Measure the accuracy
kpi_ML(Y_train, Y_train_pred, Y_test, Y_test_pred, name='ETRx200')
```

[11] ✓ 4.7s

	MAE	RMSE	Bias
ETRx200			
Train	14.1	35.6	-0.0
Test	17.6	44.9	2.4

```
# ETR Optimized x 200 Forecast for the next month
df = import_data()
X_train, Y_train, X_test, Y_test = datasets(df, x_len=12, y_len=1, test_loops=0)

ETR = ExtraTreesRegressor(n_jobs=-1, n_estimators=200, **ETR_cv.best_params_)
ETR = ETR.fit(X_train, Y_train)
forecast = pd.DataFrame(data=ETR.predict(X_test), index=df.index)
print('ETR Optimized x 200 Forecast for the next period')
display(forecast)
```

[12] ✓ 4.9s

ETR Optimized x 200 Forecast for the next period

Product	0
Product 1	1508.069742
Product 10	809.593151
Product 11	1085.162544
Product 12	170.495017
Product 13	288.021175
...	...
Product 65	0.233671
Product 66	8.463769
Product 7	284.940312
Product 8	663.120976
Product 9	141.181486

66 rows × 1 columns



# MACHINE LEARNING

## 5 ADAPTATIVE BOOSTING VS OPTIMAL ADAPTATIVE BOOSTING

DATA SCIENCE FOR SUPPLY CHAIN  
FORECASTING

```
from sklearn.tree import DecisionTreeRegressor
from sklearn.ensemble import AdaBoostRegressor
ada = AdaBoostRegressor(DecisionTreeRegressor(max_depth=8), n_estimators=100, learning_rate=0.25, loss='square')

# Fit the AdaBoost to the training data
ada = ada.fit(X_train, Y_train)

# Create predictions on the training and test sets
Y_train_pred = ada.predict(X_train)
Y_test_pred = ada.predict(X_test)

# Measure the accuracy
kpi_ML(Y_train, Y_train_pred, Y_test, Y_test_pred, name='AdaBoost')
```

[15] ✓ 14.8s

...	MAE	RMSE	Bias
AdaBoost			
Train	9.9	21.0	-0.4
Test	18.0	47.7	2.7

```
# n_estimators = [100]
learning_rate = [0.005, 0.01, 0.05, 0.1, 0.15, 0.2, 0.25, 0.3, 0.35]
loss = ['square', 'exponential', 'linear']
param_dist = {'n_estimators': n_estimators,
              'learning_rate': learning_rate,
              'loss': loss}

def model_mae(model, X, Y):
    Y_pred = model.predict(X)
    mae = np.mean(np.abs(Y - Y_pred))/np.mean(Y)
    return mae

# Let's now go into our optimization loop.
# for each max_depth tried by using ada_cv.best_score and ada_cv.best_params_

from sklearn.model_selection import RandomizedSearchCV
results = [] # To record the best score and parameters obtained

for max_depth in range(2, 18, 2):
    ada = AdaBoostRegressor(DecisionTreeRegressor(max_depth=max_depth))
    ada_cv = RandomizedSearchCV(ada, param_dist, n_jobs=-1, cv=6, n_iter=20, scoring='neg_mean_absolute_error')
    ada_cv.fit(X_train, Y_train)
    print('Tuned AdaBoost Parameters:', ada_cv.best_params_)
    print('Result:', ada_cv.best_score_)
    results.append([ada_cv.best_score_, ada_cv.best_params_, max_depth])

# We can then transform results into a DataFrame
# The method idxmax() on our DataFrame to print the parameter set that got the lowest error
# RandomizedSearchCV is returning the negative mean absolute error.

results = pd.DataFrame(data=results, columns=['Score', 'Best Params', 'Max_Depth'])
optimal = results['Score'].idxmax()
print(results.iloc[optimal])
```

[17] ✓ 51m 4.4s

Best Params	{'loss': 'exponential', 'learning_rate': 0.01}
Max_Depth	12



# MACHINE LEARNING

## 5 OPTIMAL ADAPTATIVE BOOSTING

### DATA SCIENCE FOR SUPPLY CHAIN FORECASTING

```
# Test the optimal parameter found against our training set
ada = AdaBoostRegressor(DecisionTreeRegressor(max_depth=12),
                        n_estimators=100,
                        learning_rate=0.01,
                        loss='exponential')

# Fit the AdaBoost to the training data
ada.fit(X_train, Y_train)

# Create predictions for the training and test sets
Y_train_pred = ada.predict(X_train)
Y_test_pred = ada.predict(X_test)

# Measure the accuracy
kpi ML(Y_train, Y_train_pred, Y_test, Y_test_pred, name='AdaBoost optimized')
```

[18] ✓ 17.5s

	MAE	RMSE	Bias
AdaBoost optimized			
Train	3.4	9.9	-0.0
Test	17.8	48.2	4.2

```
# AdaBoost Optimized Forecast for the next month
df = import_data()
X_train, Y_train, X_test, Y_test = datasets(df, x_len=12, y_len=1, test_loops=0)

ada = AdaBoostRegressor(DecisionTreeRegressor(max_depth=12),
                        n_estimators=100,
                        learning_rate=0.01,
                        loss='exponential')

ada = ada.fit(X_train, Y_train)
forecast = pd.DataFrame(data=ada.predict(X_test), index=df.index)
print('AdaBoost Optimized Forecast for the next period')
display(forecast)
```

[19] ✓ 18.5s

AdaBoost Optimized Forecast for the next period

Product	0
Product 1	1547.000000
Product 10	814.233333
Product 11	1287.000000
Product 12	168.816901
Product 13	305.625000
...	...
Product 65	0.060267
Product 66	9.000000
Product 7	288.857143
Product 8	658.487603
Product 9	118.000000

66 rows × 1 columns

# WRAP-UP

## BEST OPTIMAL MACHINE LEARNING MODEL FOR OUR PRODUCTS

DATA SCIENCE FOR SUPPLY CHAIN  
FORECASTING

### FOR MAE (TEST SET)

...	MAE	RMSE	Bias
Forestx200			
Train	11.9	30.4	-0.0
Test	17.5	45.4	2.3

... Forest Optimized x 200 Forecast for the next period

</>	0
Product	
Product 1	1525.167175
Product 10	804.875529
Product 11	1078.683406
Product 12	169.169076
Product 13	305.843355
...	...
Product 65	0.086483
Product 66	9.101335
Product 7	295.069647
Product 8	692.907534
Product 9	121.680588

66 rows x 1 columns

### FOR RMSE (TEST SET)

...	MAE	RMSE	Bias
ETRx200			
Train	14.1	35.6	-0.0
Test	17.6	44.9	2.4

... ETR Optimized x 200 Forecast for the next period

</>	0
Product	
Product 1	1508.069742
Product 10	809.593151
Product 11	1085.162544
Product 12	170.495017
Product 13	288.021175
...	...
Product 65	0.233671
Product 66	8.463769
Product 7	284.940312
Product 8	663.120976
Product 9	141.181486

66 rows x 1 columns

+

o

•

# THANKS.



Diego Beteta

[linkedin.com/in/diego-beteta/](https://linkedin.com/in/diego-beteta/)